

Computational Intelligence Lab

Assignment 7

Karolina Kotlowska IO czw. 9:30

7.1 Load and preprocess IMDB reviews

7.1.1 Load IMDB reviews dataset

```
In [1]: import tensorflow_datasets as tfds
import tensorflow as tf

ds_train = tfds.load('imdb_reviews', split='train', as_supervised=True, shuffle_files=True)
ds_test = tfds.load('imdb_reviews', split='test', as_supervised=True, shuffle_files=True)
```

/Library/Frameworks/Python.framework/Versions/3.10/lib/python3.10/site-packages/tqdm/auto.py:21: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html

from .autonotebook import tqdm as notebook_tqdm

```
In [2]: import pandas as pd
data = [(text.numpy().decode('UTF8'), label.numpy()) for text, label in ds_train]
df_train = pd.DataFrame(data, columns=['text', 'label'])
df_train.head()
```

2023-04-20 17:46:59.998542: W tensorflow/tsl/platform/profile_utils/cpu_utils.cc:128] Failed to get CPU frequency: 0 Hz

```
Out[2]:
```

	text	label
0	This was an absolutely terrible movie. Don't b...	0
1	I have been known to fall asleep during films,...	0
2	Mann photographs the Alberta Rocky Mountains i...	0
3	This is the kind of film for a snowy Sunday af...	1
4	As others have mentioned, all the women that g...	1

```
In [3]: data = [(text.numpy().decode('UTF8'), label.numpy()) for text, label in ds_test]
df_test = pd.DataFrame(data, columns=['text', 'label'])
df_test.head()
```

Out [3]:

	text	label
0	There are films that make careers. For George ...	1
1	A blackly comic tale of a down-trodden priest,...	1
2	Scary Movie 1-4, Epic Movie, Date Movie, Meet ...	0
3	Poor Shirley MacLaine tries hard to lend some ...	0
4	As a former Erasmus student I enjoyed this fil...	1

7.1.2 Basic preprocessing

We define `clean_text` function, which removes separators, stopwords and optionally applies lemmatization (conversion of words to basic forms)

```
In [4]: import re
from nltk.corpus import stopwords
import nltk
nltk.download('stopwords')
stopwords = stopwords.words('english')
nltk.download('wordnet')

from nltk.stem import WordNetLemmatizer

def clean_text(text, stopwords, lemmatize=False):
    text = text.lower()
    text = re.sub(r"<br />", "", text)
    # text = re.sub(r"[,;`^\.\"'!?:_()%&{}*+|#\$/-/\|<>=@|~\]\[\]", " ", text)
    text = re.sub(r"[,;`^\.\"'!?:_()%&{}*+|#\$/-/\|<>=@|~\]\[\]", " ", text)
    words = text.split()
    words = list(filter(lambda w: not w in stopwords, words))
    if lemmatize:
        lemmatizer = WordNetLemmatizer()
        words = [lemmatizer.lemmatize(token) for token in words]
        words = [lemmatizer.lemmatize(token, "v") for token in words]
    return ' '.join(words)
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] /Users/spoton/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /Users/spoton/nltk_data...
```

Small test

```
In [5]: clean_text("he's running for a president", stopwords)
```

```
Out[5]: 'running president'
```

```
In [6]: df_train['text_preprocessed'] = df_train.text.apply(lambda t: clean_text(t, stopwords))
df_test['text_preprocessed'] = df_test.text.apply(lambda t: clean_text(t, stopwords))
```

```
In [7]: df_train.head()
```

Out [7]:

	text	label	text_preprocessed
0	This was an absolutely terrible movie. Don't b...	0	absolutely terrible movie lured christopher wa...
1	I have been known to fall asleep during films,...	0	known fall asleep films usually due combinatio...
2	Mann photographs the Alberta Rocky Mountains i...	0	mann photographs alberta rocky mountains super...
3	This is the kind of film for a snowy Sunday af...	1	kind film snowy sunday afternoon rest world go...
4	As others have mentioned, all the women that g...	1	others mentioned women go nude film mostly abs...

7.1.3 Conversion to sequences

In further experiments we will use `text_preprocessed` column

```
In [8]: from keras.preprocessing.text import Tokenizer
tokenizer = Tokenizer(num_words=10_000)
tokenizer.fit_on_texts(df_train.text_preprocessed)
seq_train = tokenizer.texts_to_sequences(df_train.text_preprocessed)
seq_test = tokenizer.texts_to_sequences(df_test.text_preprocessed)
```

```
In [9]: # from keras import preprocessing
# maxlen = 500
# seq_train_padded = preprocessing.sequence.pad_sequences(seq_train, maxlen)
# seq_test_padded = preprocessing.sequence.pad_sequences(seq_test, maxlen)

# FIXME

# from keras import preprocessing
import keras
maxlen = 500
seq_train_padded = keras.utils.pad_sequences(seq_train, maxlen=maxlen)
seq_test_padded = keras.utils.pad_sequences(seq_test, maxlen=maxlen)
```

7.2 Simple RNN

Simple RNN accepts a sequence of inputs $[x(0), x(1), \dots, x(maxlen - 1)]$ and applies the following formula for updating internal state h :

$$h(t + 1) = f(W \cdot x(t) + U \cdot h(t) + b)$$

- $x(t)$ - input sequence $[x(0), x(1), \dots, x(maxlen - 1)]$
- $h(t)$ hidden state passed between iterations

Finally, at the end of sequence it yields $h(maxlen - 1)$

```
In [10]: from keras.models import Sequential
from keras.layers import Flatten, Dense
from keras.layers import Embedding, SimpleRNN
from keras.layers import GlobalMaxPool1D, MaxPool1D, Dropout
```

```

batch_size = 256
max_features=10_000

model = Sequential()
model.add(Embedding(input_dim=10_000, output_dim=64, input_length=maxlen))
model.add(SimpleRNN(units = 32, return_sequences=True))
model.add(GlobalMaxPool1D())
model.add(Dense(20, activation="relu"))
model.add(Dropout(0.05))
model.add(Dense(2, activation='softmax'))
# model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='sparse_categorical_crossentropy')
# model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 500, 64)	640000
simple_rnn (SimpleRNN)	(None, 500, 32)	3104
global_max_pooling1d (GlobalMaxPooling1D)	(None, 32)	0
dense (Dense)	(None, 20)	660
dropout (Dropout)	(None, 20)	0
dense_1 (Dense)	(None, 2)	42
Total params: 643,806		
Trainable params: 643,806		
Non-trainable params: 0		

In [11]:

```

epochs = 5
batch_size = 512
hist = model.fit(seq_train_padded, df_train.label, epochs = epochs, batch

```

```

Epoch 1/5
40/40 [=====] - 13s 320ms/step - loss: 0.6880 -
accuracy: 0.5361 - val_loss: 0.6788 - val_accuracy: 0.5262
Epoch 2/5
40/40 [=====] - 13s 334ms/step - loss: 0.5842 -
accuracy: 0.7293 - val_loss: 0.6368 - val_accuracy: 0.6134
Epoch 3/5
40/40 [=====] - 13s 333ms/step - loss: 0.4302 -
accuracy: 0.8201 - val_loss: 0.3977 - val_accuracy: 0.8256
Epoch 4/5
40/40 [=====] - 14s 347ms/step - loss: 0.3299 -
accuracy: 0.8711 - val_loss: 0.3461 - val_accuracy: 0.8518
Epoch 5/5
40/40 [=====] - 14s 343ms/step - loss: 0.2617 -
accuracy: 0.9015 - val_loss: 0.3985 - val_accuracy: 0.8150

```

TODO 7.2.1 Predict probabilities, then labels and print the classification report

```
In [12]: from sklearn.metrics import confusion_matrix, classification_report
import numpy as np

y_test = [label.numpy() for img, label in ds_test]
y_test = np.array(y_test)
probs = model.predict(seq_test_padded)
y_pred = np.argmax(probs, axis=1)

print(classification_report(y_test, y_pred))
```

```
782/782 [=====] - 10s 12ms/step
              precision    recall  f1-score   support

         0         0.93        0.67        0.77        12500
         1         0.74        0.95        0.83        12500


 accuracy          0.81        25000
 macro avg         0.83        0.81        0.80        25000
 weighted avg         0.83        0.81        0.80        25000
```

TODO 7.2.2 Plot the data collected in the history

```
In [14]: import matplotlib.pyplot as plt

plt.plot(hist.history['loss'], label='loss')
plt.plot(hist.history['val_loss'], label='val_loss')
plt.legend()
```


Out[14]: <matplotlib.legend.Legend at 0x2cd4e6fb0>



```
In [15]: plt.plot(hist.history['accuracy'], label='acc')
plt.plot(hist.history['val_accuracy'], label='val_acc')
```

```
plt.legend()
```

Out[15]: <matplotlib.legend.Legend at 0x2d7f2e3b0>



7.3 Other recurrent layers implemented in keras

- LSTM
- Bidirectional(LSTM)
- GRU

Long-Term Short Term Memory LSTM networks

You may read about LSTM networks You may read about LSTM networks [in this article](#).

LSTM is a special kind of RNN which is mainly useful for learning long-term dependencies. The name refers to the idea that the activations of a network correspond to short-term memory, while the weights correspond to long-term memory. If the activations can preserve information over long distances, that makes them long-term short-term memory.

Bidirectional

The idea of Bidirectional Recurrent Neural Networks (RNNs) is straightforward. It involves duplicating the first recurrent layer in the network so that there are now two layers side-by-side, then providing the input sequence as-is as input to the first layer and providing a reversed copy of the input sequence to the second.

[Open the article...](#)

Gated Recurrent Unit - GRU Citing [Wikipedia](#)

Gated recurrent units (GRUs) are a gating mechanism in recurrent neural networks, introduced in 2014 by Kyunghyun Cho et al. The GRU is like a long short-term memory (LSTM) with a forget gate, but has fewer parameters than LSTM, as it lacks an output gate. GRU's performance on certain tasks of polyphonic music modeling, speech signal modeling and natural language processing was found to be similar to that of LSTM. GRUs have been shown to exhibit better performance on certain smaller and less frequent datasets

We will define a few sample configurations and a function

`build_recurrent_model` that will build a model inserting a recurrent layer passed as a parameter

```
In [19]: from keras.layers import LSTM
from keras.layers import Bidirectional
from keras.layers import GRU

configurations=[
    {'recurrent_layer':SimpleRNN(units = 20,return_sequences=True),'embedding_dim':1000},
    {'recurrent_layer':LSTM(units = 20,return_sequences=True),'embedding_dim':1000},
    {'recurrent_layer':Bidirectional(LSTM(units = 8,return_sequences=True),'embedding_dim':1000)},
    {'recurrent_layer':GRU(units = 40,return_sequences=True),'embedding_dim':1000}
]

def build_recurrent_model(embedding_dim, recurrent_layer):
    model= Sequential()
    model.add(Embedding(input_dim=10_000, output_dim=embedding_dim, input_length=1000))
    model.add(recurrent_layer)
    model.add(GlobalMaxPool1D())
    model.add(Dense(20, activation="relu"))
    model.add(Dropout(0.05))
    model.add(Dense(2, activation='softmax'))
    # model.add(Dense(1, activation='sigmoid'))
    model.compile(optimizer='rmsprop', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
    # model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])
    model.summary()
    return model
```

TODO 7.3.1 Prepare at least four similar configurations using `SimpleRNN`, `LSTM`, `Bidirectional` and `GRU` layers. Each of basic configurations should be tested (or cross-validated). As a rule - you may use a configuration with smaller number of epochs for CV (cross-validation) and copy it entering greater number of epochs (for testing).

7.3.1 Using scikit learn compatible keras wrapper

We will try `KerasClassifier`, which can be used within scikit-learn pipelines and functions for hyperparameter selection (e.g. `GridSearchCV`)

```
In [17]: from keras.wrappers.scikit_learn import KerasClassifier
        #or
        #!/pip install scikeras tensorflow
        #from scikeras.wrappers import KerasClassifier
```

```
In [20]: config = configurations[2]

classifier = KerasClassifier(build_fn = lambda : build_recurrent_model(
    embedding_dim=config['embedding_dim'],
    recurrent_layer=config['recurrent_layer']),
    epochs=config['epochs'],
    batch_size=256)

classifier.fit(seq_train_padded,df_train.label)
y_pred = classifier.predict(seq_test_padded)
print(classification_report(df_test.label,y_pred))
```

```
/var/folders/db/mk99b90s0f95rb60mwwyfc9m0000gq/T/ipykernel_84877/3045942
352.py:3: DeprecationWarning: KerasClassifier is deprecated, use Sci-Ker
as (https://github.com/adriangb/scikeras) instead. See https://www.adria
ngb.com/scikeras/stable/migration.html for help migrating.
```

```
    classifier = KerasClassifier(build_fn = lambda : build_recurrent_model
(
```


Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 500, 100)	1000000
bidirectional_1 (Bidirectional)	(None, 500, 16)	6976
global_max_pooling1d_2 (GlobalMaxPooling1D)	(None, 16)	0
dense_4 (Dense)	(None, 20)	340
dropout_2 (Dropout)	(None, 20)	0
dense_5 (Dense)	(None, 2)	42

=====
 Total params: 1,007,358
 Trainable params: 1,007,358
 Non-trainable params: 0

```
Epoch 1/3
98/98 [=====] - 28s 262ms/step - loss: 0.6114 - accuracy: 0.6996
Epoch 2/3
98/98 [=====] - 26s 261ms/step - loss: 0.3088 - accuracy: 0.8768
Epoch 3/3
98/98 [=====] - 26s 264ms/step - loss: 0.2218 - accuracy: 0.9154
782/782 [=====] - 13s 16ms/step
```

	precision	recall	f1-score	support
0	0.88	0.88	0.88	12500
1	0.88	0.88	0.88	12500
accuracy			0.88	25000
macro avg	0.88	0.88	0.88	25000
weighted avg	0.88	0.88	0.88	25000

7.3.2 Applying cross-validation

We will demonstrate cross-validation using `cross_validate` function.

TODO 7.3.2 Apply CV for 4 basic configurations with a small number of epochs (e.g.3) and number of folds (also 3). Use training dataset in cross validation.

```
In [21]: import sklearn
sorted(sklearn.metrics.SCORERS.keys())
```

```
Out[21]: ['accuracy',
          'adjusted_mutual_info_score',
          'adjusted_rand_score',
          'average_precision',
          'balanced_accuracy',
          'completeness_score',
          'explained_variance',
          'f1',
          'f1_macro',
          'f1_micro',
          'f1_samples',
          'f1_weighted',
          'fowlkes_mallows_score',
          'homogeneity_score',
          'jaccard',
          'jaccard_macro',
          'jaccard_micro',
          'jaccard_samples',
          'jaccard_weighted',
          'matthews_corrcoef',
          'max_error',
          'mutual_info_score',
          'neg_brier_score',
          'neg_log_loss',
          'neg_mean_absolute_error',
          'neg_mean_absolute_percentage_error',
          'neg_mean_gamma_deviance',
          'neg_mean_poisson_deviance',
          'neg_mean_squared_error',
          'neg_mean_squared_log_error',
          'neg_median_absolute_error',
          'neg_negative_likelihood_ratio',
          'neg_root_mean_squared_error',
          'normalized_mutual_info_score',
          'positive_likelihood_ratio',
          'precision',
          'precision_macro',
          'precision_micro',
          'precision_samples',
          'precision_weighted',
          'r2',
          'rand_score',
          'recall',
          'recall_macro',
          'recall_micro',
          'recall_samples',
          'recall_weighted',
          'roc_auc',
          'roc_auc_ovo',
          'roc_auc_ovo_weighted',
          'roc_auc_ovr',
          'roc_auc_ovr_weighted',
          'top_k_accuracy',
          'v_measure_score']
```

```
In [107... from sklearn.model_selection import cross_validate

config = configurations[0]

classifier = KerasClassifier(build_fn = lambda : build_recurrent_model(
```

```

embedding_dim=config['embedding_dim'],
recurrent_layer=config['recurrent_layer']),
epochs=config['epochs'],
batch_size=256)

n_cv_folds = 3
scoring = ['accuracy', 'precision_macro', 'recall_macro', 'f1_macro', 'roc_auc']
results = cross_validate(classifier, seq_train_padded, df_train.label,
                        cv=n_cv_folds, scoring=scoring, return_train_score=False)

```

Model: "sequential_28"

Layer (type)	Output Shape	Param #
embedding_21 (Embedding)	(None, 500, 8)	80000
simple_rnn_2 (SimpleRNN)	(None, 500, 20)	580
global_max_pooling1d_21 (GlobalMaxPooling1D)	(None, 20)	0
dense_55 (Dense)	(None, 20)	420
dropout_19 (Dropout)	(None, 20)	0
dense_56 (Dense)	(None, 2)	42

```

Total params: 81,042
Trainable params: 81,042
Non-trainable params: 0

```

Epoch 1/10

```

/var/folders/db/mk99b90s0f95rb60mwwyfc9m0000gq/T/ipykernel_84877/3030849
369.py:5: DeprecationWarning: KerasClassifier is deprecated, use Sci-Ker
as (https://github.com/adriangb/scikeras) instead. See https://www.adria
ngb.com/scikeras/stable/migration.html for help migrating.
  classifier = KerasClassifier(build_fn = lambda : build_recurrent_model
(

```

```

66/66 [=====] - 4s 52ms/step - loss: 0.6093 - a
ccuracy: 0.6783
Epoch 2/10
66/66 [=====] - 3s 52ms/step - loss: 0.3962 - a
ccuracy: 0.8429
Epoch 3/10
66/66 [=====] - 3s 52ms/step - loss: 0.2970 - a
ccuracy: 0.8868
Epoch 4/10
66/66 [=====] - 3s 52ms/step - loss: 0.2338 - a
ccuracy: 0.9123
Epoch 5/10
66/66 [=====] - 3s 52ms/step - loss: 0.1888 - a
ccuracy: 0.9317
Epoch 6/10
66/66 [=====] - 3s 52ms/step - loss: 0.1578 - a
ccuracy: 0.9459
Epoch 7/10
66/66 [=====] - 3s 52ms/step - loss: 0.1299 - a
ccuracy: 0.9544
Epoch 8/10
66/66 [=====] - 3s 52ms/step - loss: 0.1113 - a
ccuracy: 0.9618
Epoch 9/10
66/66 [=====] - 3s 52ms/step - loss: 0.0911 - a
ccuracy: 0.9702
Epoch 10/10
66/66 [=====] - 3s 52ms/step - loss: 0.0823 - a
ccuracy: 0.9745
261/261 [=====] - 2s 6ms/step
261/261 [=====] - 2s 6ms/step
521/521 [=====] - 3s 6ms/step
521/521 [=====] - 3s 6ms/step
Model: "sequential_29"

```

Layer (type)	Output Shape	Param #
embedding_22 (Embedding)	(None, 500, 8)	80000
simple_rnn_2 (SimpleRNN)	(None, 500, 20)	580
global_max_pooling1d_22 (GlobalMaxPooling1D)	(None, 20)	0
dense_57 (Dense)	(None, 20)	420
dropout_20 (Dropout)	(None, 20)	0
dense_58 (Dense)	(None, 2)	42

```

=====
Total params: 81,042
Trainable params: 81,042
Non-trainable params: 0

```

```

Epoch 1/10
66/66 [=====] - 5s 54ms/step - loss: 0.6253 - a
ccuracy: 0.6460
Epoch 2/10
66/66 [=====] - 4s 54ms/step - loss: 0.4167 - a

```

```

ccuracy: 0.8318
Epoch 3/10
66/66 [=====] - 4s 53ms/step - loss: 0.3015 - a
ccuracy: 0.8838
Epoch 4/10
66/66 [=====] - 3s 53ms/step - loss: 0.2389 - a
ccuracy: 0.9110
Epoch 5/10
66/66 [=====] - 3s 53ms/step - loss: 0.1872 - a
ccuracy: 0.9345
Epoch 6/10
66/66 [=====] - 4s 53ms/step - loss: 0.1565 - a
ccuracy: 0.9473
Epoch 7/10
66/66 [=====] - 3s 52ms/step - loss: 0.1321 - a
ccuracy: 0.9572
Epoch 8/10
66/66 [=====] - 3s 53ms/step - loss: 0.1135 - a
ccuracy: 0.9644
Epoch 9/10
66/66 [=====] - 4s 54ms/step - loss: 0.0920 - a
ccuracy: 0.9719
Epoch 10/10
66/66 [=====] - 3s 53ms/step - loss: 0.0798 - a
ccuracy: 0.9765
261/261 [=====] - 2s 6ms/step
261/261 [=====] - 2s 6ms/step
521/521 [=====] - 3s 6ms/step
521/521 [=====] - 4s 7ms/step
Model: "sequential_30"

```

Layer (type)	Output Shape	Param #
embedding_23 (Embedding)	(None, 500, 8)	80000
simple_rnn_2 (SimpleRNN)	(None, 500, 20)	580
global_max_pooling1d_23 (GlobalMaxPooling1D)	(None, 20)	0
dense_59 (Dense)	(None, 20)	420
dropout_21 (Dropout)	(None, 20)	0
dense_60 (Dense)	(None, 2)	42

```

=====
Total params: 81,042
Trainable params: 81,042
Non-trainable params: 0

```

```

Epoch 1/10
66/66 [=====] - 4s 53ms/step - loss: 0.5661 - a
ccuracy: 0.7022
Epoch 2/10
66/66 [=====] - 4s 54ms/step - loss: 0.3717 - a
ccuracy: 0.8477
Epoch 3/10
66/66 [=====] - 4s 55ms/step - loss: 0.2646 - a
ccuracy: 0.9008

```

```

Epoch 4/10
66/66 [=====] - 4s 53ms/step - loss: 0.2073 - a
ccuracy: 0.9263
Epoch 5/10
66/66 [=====] - 4s 54ms/step - loss: 0.1671 - a
ccuracy: 0.9427
Epoch 6/10
66/66 [=====] - 4s 54ms/step - loss: 0.1304 - a
ccuracy: 0.9563
Epoch 7/10
66/66 [=====] - 3s 52ms/step - loss: 0.1050 - a
ccuracy: 0.9663
Epoch 8/10
66/66 [=====] - 3s 52ms/step - loss: 0.0840 - a
ccuracy: 0.9738
Epoch 9/10
66/66 [=====] - 3s 52ms/step - loss: 0.0723 - a
ccuracy: 0.9776
Epoch 10/10
66/66 [=====] - 3s 52ms/step - loss: 0.0554 - a
ccuracy: 0.9842
261/261 [=====] - 2s 7ms/step
261/261 [=====] - 2s 6ms/step
521/521 [=====] - 4s 7ms/step
521/521 [=====] - 3s 6ms/step
fit_time                35.691582
score_time              3.534172
test_accuracy           0.855040
train_accuracy          0.991620
test_precision_macro    0.855443
train_precision_macro   0.991628
test_recall_macro       0.855106
train_recall_macro      0.991617
test_f1_macro           0.855011
train_f1_macro           0.991620
test_roc_auc            0.931138
train_roc_auc           0.997963
dtype: float64
-----

```

```

In [109.. df_results = pd.DataFrame(results)
df_mean = df_results.mean()
print(df_mean)

```

```

fit_time                35.691582
score_time              3.534172
test_accuracy           0.855040
train_accuracy          0.991620
test_precision_macro    0.855443
train_precision_macro   0.991628
test_recall_macro       0.855106
train_recall_macro      0.991617
test_f1_macro           0.855011
train_f1_macro           0.991620
test_roc_auc            0.931138
train_roc_auc           0.997963
dtype: float64

```

```
In [110... config = configurations[1]

classifier = KerasClassifier(build_fn = lambda : build_recurrent_model(
    embedding_dim=config['embedding_dim'],
    recurrent_layer=config['recurrent_layer']),
    epochs=config['epochs'],
    batch_size=256)

n_cv_folds = 3
scoring = ['accuracy', 'precision_macro', 'recall_macro', 'f1_macro', 'roc_auc']
results = cross_validate(classifier, seq_train_padded, df_train.label,
                        cv=n_cv_folds, scoring=scoring, return_train_score=False)
```

Model: "sequential_31"

Layer (type)	Output Shape	Param #
embedding_24 (Embedding)	(None, 500, 100)	10000000
lstm_2 (LSTM)	(None, 500, 20)	9680
global_max_pooling1d_24 (GlobalMaxPooling1D)	(None, 20)	0
dense_61 (Dense)	(None, 20)	420
dropout_22 (Dropout)	(None, 20)	0
dense_62 (Dense)	(None, 2)	42
Total params: 1,010,142		
Trainable params: 1,010,142		
Non-trainable params: 0		

```
/var/folders/db/mk99b90s0f95rb60mwwyfc9m0000gq/T/ipykernel_84877/3799207936.py:3: DeprecationWarning: KerasClassifier is deprecated, use Sci-Keras (https://github.com/adriangb/scikeras) instead. See https://www.adriangb.com/scikeras/stable/migration.html for help migrating.
  classifier = KerasClassifier(build_fn = lambda : build_recurrent_model(
```

```

Epoch 1/10
66/66 [=====] - 22s 318ms/step - loss: 0.5017 -
accuracy: 0.7629
Epoch 2/10
66/66 [=====] - 19s 291ms/step - loss: 0.2857 -
accuracy: 0.8890
Epoch 3/10
66/66 [=====] - 19s 292ms/step - loss: 0.2119 -
accuracy: 0.9210
Epoch 4/10
66/66 [=====] - 19s 293ms/step - loss: 0.1631 -
accuracy: 0.9429
Epoch 5/10
66/66 [=====] - 19s 294ms/step - loss: 0.1263 -
accuracy: 0.9590
Epoch 6/10
66/66 [=====] - 20s 301ms/step - loss: 0.1018 -
accuracy: 0.9677
Epoch 7/10
66/66 [=====] - 19s 287ms/step - loss: 0.0812 -
accuracy: 0.9755
Epoch 8/10
66/66 [=====] - 19s 286ms/step - loss: 0.0681 -
accuracy: 0.9800
Epoch 9/10
66/66 [=====] - 20s 296ms/step - loss: 0.0538 -
accuracy: 0.9849
Epoch 10/10
66/66 [=====] - 20s 298ms/step - loss: 0.0417 -
accuracy: 0.9881
261/261 [=====] - 4s 16ms/step
261/261 [=====] - 4s 16ms/step
521/521 [=====] - 9s 17ms/step
521/521 [=====] - 8s 16ms/step
Model: "sequential_32"

```

Layer (type)	Output Shape	Param #
embedding_25 (Embedding)	(None, 500, 100)	1000000
lstm_2 (LSTM)	(None, 500, 20)	9680
global_max_pooling1d_25 (GlobalMaxPooling1D)	(None, 20)	0
dense_63 (Dense)	(None, 20)	420
dropout_23 (Dropout)	(None, 20)	0
dense_64 (Dense)	(None, 2)	42

```

=====
Total params: 1,010,142
Trainable params: 1,010,142
Non-trainable params: 0

```

```

Epoch 1/10
66/66 [=====] - 22s 308ms/step - loss: 0.6366 -
accuracy: 0.6226
Epoch 2/10

```



```

66/66 [=====] - 20s 303ms/step - loss: 0.3830 -
accuracy: 0.8514
Epoch 3/10
66/66 [=====] - 20s 302ms/step - loss: 0.2526 -
accuracy: 0.9067
Epoch 4/10
66/66 [=====] - 20s 304ms/step - loss: 0.1834 -
accuracy: 0.9363
Epoch 5/10
66/66 [=====] - 20s 306ms/step - loss: 0.1372 -
accuracy: 0.9540
Epoch 6/10
66/66 [=====] - 19s 292ms/step - loss: 0.1091 -
accuracy: 0.9651
Epoch 7/10
66/66 [=====] - 20s 306ms/step - loss: 0.0839 -
accuracy: 0.9747
Epoch 8/10
66/66 [=====] - 21s 320ms/step - loss: 0.0659 -
accuracy: 0.9810
Epoch 9/10
66/66 [=====] - 21s 314ms/step - loss: 0.0499 -
accuracy: 0.9860
Epoch 10/10
66/66 [=====] - 20s 311ms/step - loss: 0.0419 -
accuracy: 0.9888
261/261 [=====] - 5s 17ms/step
261/261 [=====] - 4s 17ms/step
521/521 [=====] - 9s 16ms/step
521/521 [=====] - 8s 16ms/step
Model: "sequential_33"

```

Layer (type)	Output Shape	Param #
embedding_26 (Embedding)	(None, 500, 100)	1000000
lstm_2 (LSTM)	(None, 500, 20)	9680
global_max_pooling1d_26 (GlobalMaxPooling1D)	(None, 20)	0
dense_65 (Dense)	(None, 20)	420
dropout_24 (Dropout)	(None, 20)	0
dense_66 (Dense)	(None, 2)	42

```

=====
Total params: 1,010,142
Trainable params: 1,010,142
Non-trainable params: 0

```

```

Epoch 1/10
66/66 [=====] - 20s 280ms/step - loss: 0.6144 -
accuracy: 0.6685
Epoch 2/10
66/66 [=====] - 19s 287ms/step - loss: 0.4088 -
accuracy: 0.8374
Epoch 3/10
66/66 [=====] - 18s 275ms/step - loss: 0.3006 -

```

```

accuracy: 0.8865
Epoch 4/10
66/66 [=====] - 19s 293ms/step - loss: 0.2322 -
accuracy: 0.9164
Epoch 5/10
66/66 [=====] - 19s 283ms/step - loss: 0.1804 -
accuracy: 0.9378
Epoch 6/10
66/66 [=====] - 19s 288ms/step - loss: 0.1428 -
accuracy: 0.9534
Epoch 7/10
66/66 [=====] - 20s 299ms/step - loss: 0.1110 -
accuracy: 0.9641
Epoch 8/10
66/66 [=====] - 19s 289ms/step - loss: 0.0864 -
accuracy: 0.9741
Epoch 9/10
66/66 [=====] - 19s 293ms/step - loss: 0.0664 -
accuracy: 0.9803
Epoch 10/10
66/66 [=====] - 20s 303ms/step - loss: 0.0528 -
accuracy: 0.9848
261/261 [=====] - 5s 17ms/step
261/261 [=====] - 4s 17ms/step
521/521 [=====] - 9s 17ms/step
521/521 [=====] - 8s 16ms/step

```

```

In [111]: df_results = pd.DataFrame(results)
df_mean = df_results.mean()
print(df_mean)

```

```

fit_time          197.578003
score_time         8.937461
test_accuracy      0.845320
train_accuracy     0.987680
test_precision_macro 0.850614
train_precision_macro 0.987884
test_recall_macro  0.845260
train_recall_macro 0.987690
test_f1_macro      0.844686
train_f1_macro     0.987679
test_roc_auc       0.927810
train_roc_auc      0.998345
dtype: float64

```

```

In [112]: config = configurations[2]

classifier = KerasClassifier(build_fn = lambda : build_recurrent_model(
    embedding_dim=config['embedding_dim'],
    recurrent_layer=config['recurrent_layer']),
    epochs=config['epochs'],
    batch_size=256)

n_cv_folds = 3
scoring = ['accuracy', 'precision_macro', 'recall_macro', 'f1_macro', 'roc_auc']
results = cross_validate(classifier, seq_train_padded, df_train.label,
                        cv=n_cv_folds, scoring=scoring, return_train_score=False)
df_results = pd.DataFrame(results)
df_mean = df_results.mean()
print(df_mean, "\n-----\n\n")

```

```
/var/folders/db/mk99b90s0f95rb60mwwyfc9m0000gq/T/ipykernel_84877/4089919225.py:3: DeprecationWarning: KerasClassifier is deprecated, use Sci-Keras (https://github.com/adriangb/scikeras) instead. See https://www.adriangb.com/scikeras/stable/migration.html for help migrating.  
    classifier = KerasClassifier(build_fn = lambda : build_recurrent_model  
(
```

Model: "sequential_34"

Layer (type)	Output Shape	Param #
embedding_27 (Embedding)	(None, 500, 100)	1000000
bidirectional_1 (Bidirectional)	(None, 500, 16)	6976
global_max_pooling1d_27 (GlobalMaxPooling1D)	(None, 16)	0
dense_67 (Dense)	(None, 20)	340
dropout_25 (Dropout)	(None, 20)	0
dense_68 (Dense)	(None, 2)	42

=====
 Total params: 1,007,358
 Trainable params: 1,007,358
 Non-trainable params: 0

Epoch 1/3
 66/66 [=====] - 19s 243ms/step - loss: 0.5947 - accuracy: 0.7069
 Epoch 2/3
 66/66 [=====] - 16s 239ms/step - loss: 0.4637 - accuracy: 0.8233
 Epoch 3/3
 66/66 [=====] - 16s 246ms/step - loss: 0.3919 - accuracy: 0.8610
 261/261 [=====] - 4s 15ms/step
 261/261 [=====] - 4s 15ms/step
 521/521 [=====] - 8s 14ms/step
 521/521 [=====] - 7s 14ms/step

Model: "sequential_35"

Layer (type)	Output Shape	Param #
embedding_28 (Embedding)	(None, 500, 100)	1000000
bidirectional_1 (Bidirectional)	(None, 500, 16)	6976
global_max_pooling1d_28 (GlobalMaxPooling1D)	(None, 16)	0
dense_69 (Dense)	(None, 20)	340
dropout_26 (Dropout)	(None, 20)	0
dense_70 (Dense)	(None, 2)	42

=====
 Total params: 1,007,358
 Trainable params: 1,007,358
 Non-trainable params: 0

Epoch 1/3

```

66/66 [=====] - 17s 231ms/step - loss: 0.6022 -
accuracy: 0.6867
Epoch 2/3
66/66 [=====] - 16s 240ms/step - loss: 0.3798 -
accuracy: 0.8595
Epoch 3/3
66/66 [=====] - 16s 243ms/step - loss: 0.2658 -
accuracy: 0.9021
261/261 [=====] - 4s 14ms/step
261/261 [=====] - 4s 14ms/step
521/521 [=====] - 7s 14ms/step
521/521 [=====] - 7s 14ms/step
Model: "sequential_36"

```

Layer (type)	Output Shape	Param #
embedding_29 (Embedding)	(None, 500, 100)	1000000
bidirectional_1 (Bidirectional)	(None, 500, 16)	6976
global_max_pooling1d_29 (GlobalMaxPooling1D)	(None, 16)	0
dense_71 (Dense)	(None, 20)	340
dropout_27 (Dropout)	(None, 20)	0
dense_72 (Dense)	(None, 2)	42

```

=====
Total params: 1,007,358
Trainable params: 1,007,358
Non-trainable params: 0

```

```

Epoch 1/3
66/66 [=====] - 18s 246ms/step - loss: 0.5258 -
accuracy: 0.7523
Epoch 2/3
66/66 [=====] - 16s 239ms/step - loss: 0.2964 -
accuracy: 0.8940
Epoch 3/3
66/66 [=====] - 16s 239ms/step - loss: 0.2117 -
accuracy: 0.9288
261/261 [=====] - 4s 15ms/step
261/261 [=====] - 4s 14ms/step
521/521 [=====] - 7s 14ms/step
521/521 [=====] - 7s 14ms/step
fit_time          50.421065
score_time        8.113526
test_accuracy     0.822559
train_accuracy    0.876180
test_precision_macro 0.840712
train_precision_macro 0.890158
test_recall_macro 0.822043
train_recall_macro 0.876384
test_f1_macro     0.819648
train_f1_macro    0.874736
test_roc_auc      0.918331
train_roc_auc     0.960108

```

dtype: float64

```
In [113... df_results = pd.DataFrame(results)
df_mean = df_results.mean()
print(df_mean)
```

```
fit_time          50.421065
score_time        8.113526
test_accuracy     0.822559
train_accuracy    0.876180
test_precision_macro 0.840712
train_precision_macro 0.890158
test_recall_macro 0.822043
train_recall_macro 0.876384
test_f1_macro     0.819648
train_f1_macro    0.874736
test_roc_auc      0.918331
train_roc_auc     0.960108
dtype: float64
```

```
In [114... config = configurations[3]

classifier = KerasClassifier(build_fn = lambda : build_recurrent_model(
    embedding_dim=config['embedding_dim'],
    recurrent_layer=config['recurrent_layer'],
    epochs=config['epochs'],
    batch_size=256)

n_cv_folds = 3
scoring = ['accuracy', 'precision_macro', 'recall_macro', 'f1_macro', 'roc_auc']
results = cross_validate(classifier, seq_train_padded, df_train.label,
                        cv=n_cv_folds, scoring=scoring, return_train_score=False)
df_results = pd.DataFrame(results)
df_mean = df_results.mean()
print(df_mean, "\n-----\n\n")
```

```
/var/folders/db/mk99b90s0f95rb60mwwyfc9m0000gq/T/ipykernel_84877/6794986
52.py:3: DeprecationWarning: KerasClassifier is deprecated, use Sci-Keras
(https://github.com/adriangb/scikeras) instead. See https://www.adriangb.com/scikeras/stable/migration.html for help migrating.
```

```
classifier = KerasClassifier(build_fn = lambda : build_recurrent_model
(
```

Model: "sequential_37"

Layer (type)	Output Shape	Param #
embedding_30 (Embedding)	(None, 500, 100)	10000000
gru_1 (GRU)	(None, 500, 40)	17040
global_max_pooling1d_30 (GlobalMaxPooling1D)	(None, 40)	0
dense_73 (Dense)	(None, 20)	820
dropout_28 (Dropout)	(None, 20)	0
dense_74 (Dense)	(None, 2)	42

=====
 Total params: 1,017,902
 Trainable params: 1,017,902
 Non-trainable params: 0

Epoch 1/3
 66/66 [=====] - 37s 535ms/step - loss: 0.6492 - accuracy: 0.6787
 Epoch 2/3
 66/66 [=====] - 37s 557ms/step - loss: 0.4455 - accuracy: 0.8078
 Epoch 3/3
 66/66 [=====] - 35s 525ms/step - loss: 0.3024 - accuracy: 0.8756
 261/261 [=====] - 6s 21ms/step
 261/261 [=====] - 5s 21ms/step
 521/521 [=====] - 11s 21ms/step
 521/521 [=====] - 11s 21ms/step

Model: "sequential_38"

Layer (type)	Output Shape	Param #
embedding_31 (Embedding)	(None, 500, 100)	10000000
gru_1 (GRU)	(None, 500, 40)	17040
global_max_pooling1d_31 (GlobalMaxPooling1D)	(None, 40)	0
dense_75 (Dense)	(None, 20)	820
dropout_29 (Dropout)	(None, 20)	0
dense_76 (Dense)	(None, 2)	42

=====
 Total params: 1,017,902
 Trainable params: 1,017,902
 Non-trainable params: 0

Epoch 1/3
 66/66 [=====] - 35s 521ms/step - loss: 0.5785 - accuracy: 0.7061

Epoch 2/3
 66/66 [=====] - 36s 547ms/step - loss: 0.3477 - accuracy: 0.8629
 Epoch 3/3
 66/66 [=====] - 38s 570ms/step - loss: 0.2577 - accuracy: 0.9013
 261/261 [=====] - 6s 21ms/step
 261/261 [=====] - 5s 21ms/step
 521/521 [=====] - 11s 21ms/step
 521/521 [=====] - 11s 21ms/step
 Model: "sequential_39"

Layer (type)	Output Shape	Param #
embedding_32 (Embedding)	(None, 500, 100)	1000000
gru_1 (GRU)	(None, 500, 40)	17040
global_max_pooling1d_32 (GlobalMaxPooling1D)	(None, 40)	0
dense_77 (Dense)	(None, 20)	820
dropout_30 (Dropout)	(None, 20)	0
dense_78 (Dense)	(None, 2)	42

=====
 Total params: 1,017,902
 Trainable params: 1,017,902
 Non-trainable params: 0

Epoch 1/3
 66/66 [=====] - 39s 574ms/step - loss: 0.5762 - accuracy: 0.7109
 Epoch 2/3
 66/66 [=====] - 37s 563ms/step - loss: 0.3716 - accuracy: 0.8498
 Epoch 3/3
 66/66 [=====] - 37s 563ms/step - loss: 0.2781 - accuracy: 0.8942
 261/261 [=====] - 6s 21ms/step
 261/261 [=====] - 5s 21ms/step
 521/521 [=====] - 11s 21ms/step
 521/521 [=====] - 11s 21ms/step
 fit_time 110.479637
 score_time 11.182815
 test_accuracy 0.839399
 train_accuracy 0.898480
 test_precision_macro 0.856232
 train_precision_macro 0.909356
 test_recall_macro 0.839242
 train_recall_macro 0.898543
 test_f1_macro 0.837125
 train_f1_macro 0.897651
 test_roc_auc 0.939744
 train_roc_auc 0.976934
 dtype: float64


```
In [115]: df_results = pd.DataFrame(results)
df_mean = df_results.mean()
print(df_mean)
```

```
fit_time          110.479637
score_time        11.182815
test_accuracy      0.839399
train_accuracy     0.898480
test_precision_macro 0.856232
train_precision_macro 0.909356
test_recall_macro  0.839242
train_recall_macro 0.898543
test_f1_macro      0.837125
train_f1_macro     0.897651
test_roc_auc       0.939744
train_roc_auc      0.976934
dtype: float64
```

TODO 7.3.3 For each configuration tested - collect scores, compute mean values and finally prepare a table summarizing the results. Expected four rows (results for a configuration in one row).

	fit_time	score_time	test_accuracy	train_accuracy	test_precision_macro	tra
Model 0	35.691582	3.534172	0.855040	0.991620	0.855443	0.9
Model 1	197.578003	8.937461	0.845320	0.987680	0.850614	0.9
Model 2	50.421065	8.113526	0.822559	0.876180	0.840712	0.8
Model 3	110.479637	11.182815	0.839399	0.898480	0.856232	0.9

TODO 7.3.4 Select two most promising configurations for testing (based on accuracy or F1). Apply more epochs. During the testing use the test dataset. Gather the results in a table (two rows), giving:

- name of the configuration
- recurrent network type
- number of epochs
- main scores

```
In [ ]: config = configurations[0]

classifier = KerasClassifier(build_fn = lambda : build_recurrent_model(
    embedding_dim=config['embedding_dim'],
    recurrent_layer=config['recurrent_layer']),
    epochs=config['epochs'],
    batch_size=256)

n_cv_folds = 3
scoring = ['accuracy', 'precision_macro', 'recall_macro', 'f1_macro', 'roc_auc']
results = cross_validate(classifier, seq_train_padded, df_train.label,
```

```

cv=n_cv_folds, scoring=scoring, return_train_score=True)
df_results = pd.DataFrame(results)
df_mean = df_results.mean()
print(df_mean, "\n-----\n\n")

```

```

In [ ]: config = configurations[3]

classifier = KerasClassifier(build_fn = lambda : build_recurrent_model(
    embedding_dim=config['embedding_dim'],
    recurrent_layer=config['recurrent_layer'],
    epochs=config['epochs'],
    batch_size=256)

n_cv_folds = 3
scoring = ['accuracy', 'precision_macro', 'recall_macro', 'f1_macro', 'roc_auc']
results = cross_validate(classifier, seq_train_padded, df_train.label,
                        cv=n_cv_folds, scoring=scoring, return_train_score=True)
df_results = pd.DataFrame(results)
df_mean = df_results.mean()
print(df_mean, "\n-----\n\n")

```

	number of epochs	fit_time	score_time	test_accuracy	train_accuracy	test_precision
SimpleRNN	9	35.691582	3.534172	0.855040	0.991620	0.855443
Bidirectional	3	110.479637	11.182815	0.839399	0.898480	0.856232

7.3.3 Conv1D - sequence processing based on 1D convolutions

```

In [26]: from keras.models import Sequential
from keras import layers
from tensorflow.keras.optimizers import RMSprop

model = Sequential()
model.add(layers.Embedding(input_dim=10_000, output_dim=128, input_length=seq_train_padded.shape[1]))
model.add(layers.Conv1D(filters=32, kernel_size=7, activation='relu'))
model.add(layers.MaxPooling1D(pool_size=5))
model.add(layers.Conv1D(filters=32, kernel_size=7, activation='relu'))
model.add(layers.GlobalMaxPooling1D())
model.add(layers.Dense(1))

model.summary()

model.compile(optimizer=RMSprop(learning_rate=1e-4),
              loss='binary_crossentropy',
              metrics=['accuracy'])
hist = model.fit(seq_train_padded, df_train.label, epochs=20, batch_size=128)

```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
embedding_10 (Embedding)	(None, 500, 128)	1280000
conv1d_2 (Conv1D)	(None, 494, 32)	28704
max_pooling1d_1 (MaxPooling 1D)	(None, 98, 32)	0
conv1d_3 (Conv1D)	(None, 92, 32)	7200
global_max_pooling1d_10 (GlobalMaxPooling1D)	(None, 32)	0
dense_19 (Dense)	(None, 1)	33

=====
 Total params: 1,315,937
 Trainable params: 1,315,937
 Non-trainable params: 0

Epoch 1/20
 157/157 [=====] - 10s 58ms/step - loss: 0.7279 - accuracy: 0.5273 - val_loss: 0.6888 - val_accuracy: 0.5298
 Epoch 2/20
 157/157 [=====] - 9s 57ms/step - loss: 0.6662 - accuracy: 0.6780 - val_loss: 0.6555 - val_accuracy: 0.7018
 Epoch 3/20
 157/157 [=====] - 9s 59ms/step - loss: 0.5948 - accuracy: 0.7871 - val_loss: 0.5436 - val_accuracy: 0.7800
 Epoch 4/20
 157/157 [=====] - 9s 59ms/step - loss: 0.4488 - accuracy: 0.8278 - val_loss: 0.4391 - val_accuracy: 0.8224
 Epoch 5/20
 157/157 [=====] - 10s 64ms/step - loss: 0.3699 - accuracy: 0.8661 - val_loss: 0.4193 - val_accuracy: 0.8444
 Epoch 6/20
 157/157 [=====] - 9s 57ms/step - loss: 0.3207 - accuracy: 0.8874 - val_loss: 0.4377 - val_accuracy: 0.8562
 Epoch 7/20
 157/157 [=====] - 9s 59ms/step - loss: 0.2814 - accuracy: 0.9036 - val_loss: 0.4592 - val_accuracy: 0.8632
 Epoch 8/20
 157/157 [=====] - 9s 60ms/step - loss: 0.2519 - accuracy: 0.9158 - val_loss: 0.5229 - val_accuracy: 0.8648
 Epoch 9/20
 157/157 [=====] - 9s 60ms/step - loss: 0.2276 - accuracy: 0.9260 - val_loss: 0.5098 - val_accuracy: 0.8704
 Epoch 10/20
 157/157 [=====] - 9s 60ms/step - loss: 0.2099 - accuracy: 0.9334 - val_loss: 0.5392 - val_accuracy: 0.8716
 Epoch 11/20
 157/157 [=====] - 9s 59ms/step - loss: 0.1895 - accuracy: 0.9416 - val_loss: 0.5690 - val_accuracy: 0.8716
 Epoch 12/20
 157/157 [=====] - 10s 61ms/step - loss: 0.1740 - accuracy: 0.9491 - val_loss: 0.6050 - val_accuracy: 0.8708
 Epoch 13/20

```
157/157 [=====] - 9s 57ms/step - loss: 0.1592 -  
accuracy: 0.9570 - val_loss: 0.7103 - val_accuracy: 0.8642  
Epoch 14/20  
157/157 [=====] - 9s 57ms/step - loss: 0.1457 -  
accuracy: 0.9633 - val_loss: 0.6752 - val_accuracy: 0.8682  
Epoch 15/20  
157/157 [=====] - 9s 57ms/step - loss: 0.1318 -  
accuracy: 0.9692 - val_loss: 0.7176 - val_accuracy: 0.8654  
Epoch 16/20  
157/157 [=====] - 9s 57ms/step - loss: 0.1208 -  
accuracy: 0.9758 - val_loss: 0.7512 - val_accuracy: 0.8650  
Epoch 17/20  
157/157 [=====] - 9s 57ms/step - loss: 0.1102 -  
accuracy: 0.9797 - val_loss: 0.8072 - val_accuracy: 0.8650  
Epoch 18/20  
157/157 [=====] - 9s 58ms/step - loss: 0.1008 -  
accuracy: 0.9844 - val_loss: 0.8529 - val_accuracy: 0.8598  
Epoch 19/20  
157/157 [=====] - 9s 57ms/step - loss: 0.0927 -  
accuracy: 0.9882 - val_loss: 0.9020 - val_accuracy: 0.8614  
Epoch 20/20  
157/157 [=====] - 9s 57ms/step - loss: 0.0862 -  
accuracy: 0.9906 - val_loss: 0.9356 - val_accuracy: 0.8584
```

TODO Compute and display scores for the test set

```
In [27]: classifier.fit(seq_train_padded,df_train.label)  
y_pred = classifier.predict(seq_test_padded)  
print(classification_report(y_pred,df_test.label))
```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
embedding_11 (Embedding)	(None, 500, 100)	10000000
bidirectional_1 (Bidirectional)	(None, 500, 16)	6976
global_max_pooling1d_11 (GlobalMaxPooling1D)	(None, 16)	0
dense_20 (Dense)	(None, 20)	340
dropout_9 (Dropout)	(None, 20)	0
dense_21 (Dense)	(None, 2)	42

=====
 Total params: 1,007,358
 Trainable params: 1,007,358
 Non-trainable params: 0

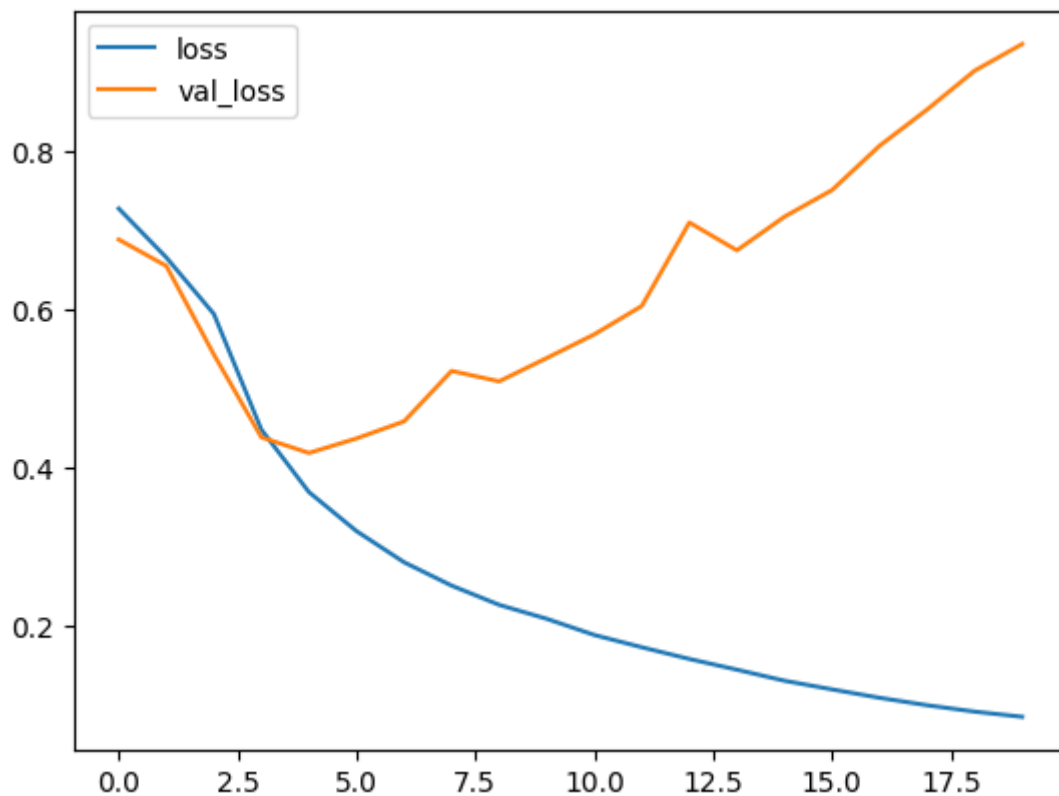
```
Epoch 1/3
98/98 [=====] - 30s 283ms/step - loss: 0.5842 - accuracy: 0.6901
Epoch 2/3
98/98 [=====] - 27s 276ms/step - loss: 0.3785 - accuracy: 0.8578
Epoch 3/3
98/98 [=====] - 27s 277ms/step - loss: 0.3377 - accuracy: 0.8860
782/782 [=====] - 13s 16ms/step
```

	precision	recall	f1-score	support
0	0.85	0.86	0.86	12329
1	0.87	0.85	0.86	12671
accuracy			0.86	25000
macro avg	0.86	0.86	0.86	25000
weighted avg	0.86	0.86	0.86	25000

TODO 7.3.5 Plot values in the history


```
In [28]: plt.plot(hist.history['loss'], label='loss')
plt.plot(hist.history['val_loss'], label='val_loss')
plt.legend()
```

```
Out[28]: <matplotlib.legend.Legend at 0x30525b640>
```



```
In [29]: plt.plot(hist.history['accuracy'], label='acc')  
plt.plot(hist.history['val_accuracy'], label='val_acc')  
plt.legend()
```

Out[29]: <matplotlib.legend.Legend at 0x2e76db1c0>



7.4 Regression - bike sharing dataset

We will use a dataset containing information on bike rentals. We will focus on the attribute `cnt` collected every hour.

```
In [61]: !wget https://dysk.agh.edu.pl/s/G6ZNziBRbEEcMeN/download -O Bike-Sharing-  
!unzip Bike-Sharing-Dataset.zip  
!cat Readme.txt
```

```
--2023-04-20 18:28:17-- https://dysk.agh.edu.pl/s/G6ZNziBRbEEcMeN/download
Resolving dysk.agh.edu.pl (dysk.agh.edu.pl)... 2001:6d8:10:1060::6004, 149.156.96.4
Connecting to dysk.agh.edu.pl (dysk.agh.edu.pl)|2001:6d8:10:1060::6004|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 279992 (273K) [application/zip]
Saving to: 'Bike-Sharing-Dataset.zip'
```

```
Bike-Sharing-Dataset 100%[=====>] 273,43K --.-KB/s in 0,05s
```

```
2023-04-20 18:28:18 (4,90 MB/s) - 'Bike-Sharing-Dataset.zip' saved [279992/279992]
```

```
Archive: Bike-Sharing-Dataset.zip
replace Readme.txt? [y]es, [n]o, [A]ll, [N]one, [r]ename: ^C
```

```
=====
Bike Sharing Dataset
=====
```

Hadi Fanaee-T

Laboratory of Artificial Intelligence and Decision Support (LIAAD), University of Porto
INESC Porto, Campus da FEUP
Rua Dr. Roberto Frias, 378
4200 - 465 Porto, Portugal

```
=====
Background
=====
```

Bike sharing systems are new generation of traditional bike rentals where the whole process from membership, rental and return back has become automatic. Through these systems, user is able to easily rent a bike from a particular position and return back at another position. Currently, there are about over 500 bike-sharing programs around the world which is composed of over 500 thousands bicycles. Today, there exists great interest in these systems due to their important role in traffic, environmental and health issues.

Apart from interesting real world applications of bike sharing systems, the characteristics of data being generated by these systems make them attractive for the research. Opposed to other transport services such as bus or subway, the duration of travel, departure and arrival position is explicitly recorded in these systems. This feature turns bike sharing system into a virtual sensor network that can be used for sensing mobility in the city. Hence, it is expected that most of important events in the city could be detected via monitoring these data.

```
=====
Data Set
=====
```

Bike-sharing rental process is highly correlated to the environmental and seasonal settings. For instance, weather conditions,

precipitation, day of week, season, hour of the day, etc. can affect the rental behaviors. The core data set is related to the two-year historical log corresponding to years 2011 and 2012 from Capital Bikeshare system, Washington D.C., USA which is publicly available in <http://capitalbikeshare.com/system-data>. We aggregated the data on two hourly and daily basis and then extracted and added the corresponding weather and seasonal information. Weather information are extracted from <http://www.freemeteo.com>.

=====

Associated tasks

=====

- Regression:

Prediction of bike rental count hourly or daily based on the environmental and seasonal settings.

- Event and Anomaly Detection:

Count of rented bikes are also correlated to some events in the town which easily are traceable via search engines.

For instance, query like "2012-10-30 washington d.c." in Google returns related results to Hurricane Sandy. Some of the important events are

identified in [1]. Therefore the data can be used for validation of anomaly or event detection algorithms as well.

=====

Files

=====

- Readme.txt

- hour.csv : bike sharing counts aggregated on hourly basis. Records: 17379 hours

- day.csv - bike sharing counts aggregated on daily basis. Records: 731 days

=====

Dataset characteristics

=====

Both hour.csv and day.csv have the following fields, except hr which is not available in day.csv

- instant: record index
- dteday : date
- season : season (1:springer, 2:summer, 3:fall, 4:winter)
- yr : year (0: 2011, 1:2012)
- mnth : month (1 to 12)
- hr : hour (0 to 23)
- holiday : weather day is holiday or not (extracted from <http://dchr.dc.gov/page/holiday-schedule>)
- weekday : day of the week
- workingday : if day is neither weekend nor holiday is 1, otherwise is 0.
- + weathersit :
 - 1: Clear, Few clouds, Partly cloudy, Partly cloudy
 - 2: Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist
 - 3: Light Snow, Light Rain + Thunderstorm + Scattered c

clouds, Light Rain + Scattered clouds
 - 4: Heavy Rain + Ice Pallets + Thunderstorm + Mist, Snow + Fog
 - temp : Normalized temperature in Celsius. The values are divided to 41 (max)
 - atemp: Normalized feeling temperature in Celsius. The values are divided to 50 (max)
 - hum: Normalized humidity. The values are divided to 100 (max)
 - windspeed: Normalized wind speed. The values are divided to 67 (max)
 - casual: count of casual users
 - registered: count of registered users
 - cnt: count of total rental bikes including both casual and registered

=====

License

=====

Use of this dataset in publications must be cited to the following publication:

[1] Fanaee-T, Hadi, and Gama, Joao, "Event labeling combining ensemble detectors and background knowledge", Progress in Artificial Intelligence (2013): pp. 1-15, Springer Berlin Heidelberg, doi:10.1007/s13748-013-0040-3.

```
@article{
  year={2013},
  issn={2192-6352},
  journal={Progress in Artificial Intelligence},
  doi={10.1007/s13748-013-0040-3},
  title={Event labeling combining ensemble detectors and background knowledge},
  url={http://dx.doi.org/10.1007/s13748-013-0040-3},
  publisher={Springer Berlin Heidelberg},
  keywords={Event labeling; Event detection; Ensemble learning; Background knowledge},
  author={Fanaee-T, Hadi and Gama, Joao},
  pages={1-15}
}
```

=====

Contact

=====

For further information about this dataset please contact Hadi Fanaee-T (hadi.fanaee@fe.up.pt)

```
In [62]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

df = pd.read_csv('hour.csv', parse_dates=['dteday'])
df.head(len(df))
```


Out [62]:

	instant	dteday	season	yr	mnth	hr	holiday	weekday	workingday	weathersi
0	1	2011-01-01	1	0	1	0	0	6	0	
1	2	2011-01-01	1	0	1	1	0	6	0	
2	3	2011-01-01	1	0	1	2	0	6	0	
3	4	2011-01-01	1	0	1	3	0	6	0	
4	5	2011-01-01	1	0	1	4	0	6	0	
...
17374	17375	2012-12-31	1	1	12	19	0	1	1	
17375	17376	2012-12-31	1	1	12	20	0	1	1	
17376	17377	2012-12-31	1	1	12	21	0	1	1	
17377	17378	2012-12-31	1	1	12	22	0	1	1	
17378	17379	2012-12-31	1	1	12	23	0	1	1	

17379 rows x 17 columns

```
In [63]: plt.rcParams["figure.figsize"] = (12,8)
plt.plot(df.cnt)
```

Out [63]: [<matplotlib.lines.Line2D at 0x2ee6fc040>]

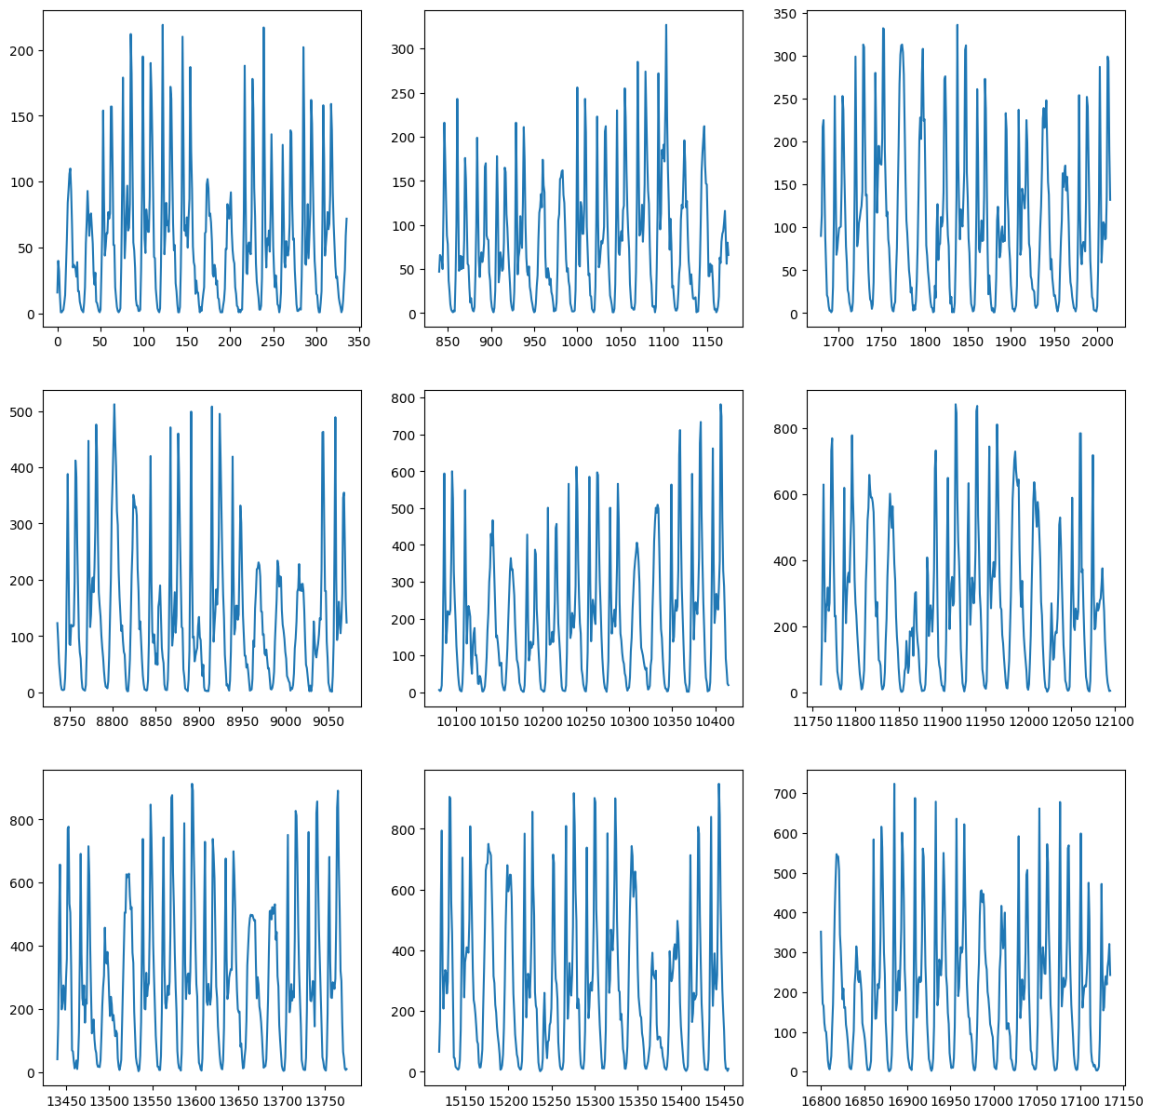


Let us plot some examples for time intervals corresponding to two weeks. Visible fluctuations with 7 days and 24 periods.

```
In [64]: import numpy as np

starts = np.array([0,5*7*24,10*7*24, 52*7*24, 60*7*24,70*7*24,80*7*24,90*
subplots = [df.cnt[starts[i]:starts[i]+14*24] for i in range(len(starts))

fig = plt.figure(figsize=(15, 15))
for i in range(9):
    ax = fig.add_subplot(3,3, 1 + i)
    ax.plot(subplots[i])
    # if title is not None:
    #     plt.suptitle(title)
plt.show()
```



Problem statement

We will use only a sequence of `cnt` values - let us denote it as $[x(0), x(1), \dots, x(n)]$.

Our idea is to predict the target value in the feature based on a sequence registered recently in the past.

Let t be a current discrete time.

- $x(t + h)$ is a target value to be predicted. $h \geq 1$ stands for *horizon*
- $[x(t - w), x(t - w + 1), \dots, x(t)]$ is a subsequence used for making predictions. It starts at $x(t - w)$ and ends at the current value $x(t)$. The parameter w stands for *window size*.

In our problem statement we select:

- h (horizon) equals to one week
- w (window size) equals to two weeks

So, we predict one week ahead.

7.4.1 Generator

The generator function will yield a tuple `(samples, targets)` where `samples` is one batch of input data and `targets` is the corresponding array of target temperatures. It takes the following arguments:

- `data` : The original array of floating point data, which we just normalized in the code snippet above.
- `window_size` : How many timesteps back should our input data go.
- `horizon` : How many timesteps in the future should our target be.
- `min_index` and `max_index` : Indices in the `data` array that delimit which timesteps to draw from. This is useful for keeping a segment of the data for validation and another one for testing.
- `shuffle` : Whether to shuffle our samples or draw them in chronological order.
- `batch_size` : The number of samples per batch.
- `step` : The period, in timesteps, at which we sample data. We will set it 6 in order to draw one data point every hour.
- `expand_dims` wrap 2D arrays as 3D tensors.

```
In [65]: import sklearn

def generator(data, window_size, horizon, min_index, max_index, shuffle=False):
    if max_index is None:
        max_index = len(data) - 1
    cursor = min_index + window_size
    while 1:
        if cursor + batch_size + horizon > max_index:
            cursor = min_index + window_size

        rows = np.arange(cursor, min(cursor + batch_size, max_index))
        cursor += len(rows)

        samples = np.zeros((len(rows), window_size // step))
        targets = np.zeros((len(rows),))
        for j, row in enumerate(rows):
            indices = np.arange(rows[j] - window_size, rows[j], step, dtype=int)
            samples[j] = data[indices]
```

```
        targets[j] = data[rows[j] + horizon-1]
    if shuffle:
        samples, targets = sklearn.utils.shuffle(samples, targets)
    if expand_dims:
        samples = np.expand_dims(samples, axis=-1)
        # targets = np.expand_dims(targets, axis=-1)
    yield samples, targets
```

Some tests for demonstration

```
In [66]: x=np.arange(3000)
gen = generator(x,10,10,min_index=0, max_index=70,batch_size = 5,shuffle=
for i,(x,y) in enumerate(gen):
    print(i,x,y)
    if i== 10: break
```

```

0 [[ 0.  1.  2.  3.  4.  5.  6.  7.  8.  9.]
   [ 1.  2.  3.  4.  5.  6.  7.  8.  9. 10.]
   [ 2.  3.  4.  5.  6.  7.  8.  9. 10. 11.]
   [ 3.  4.  5.  6.  7.  8.  9. 10. 11. 12.]
   [ 4.  5.  6.  7.  8.  9. 10. 11. 12. 13.]] [19. 20. 21. 22. 23.]
1 [[ 5.  6.  7.  8.  9. 10. 11. 12. 13. 14.]
   [ 6.  7.  8.  9. 10. 11. 12. 13. 14. 15.]
   [ 7.  8.  9. 10. 11. 12. 13. 14. 15. 16.]
   [ 8.  9. 10. 11. 12. 13. 14. 15. 16. 17.]
   [ 9. 10. 11. 12. 13. 14. 15. 16. 17. 18.]] [24. 25. 26. 27. 28.]
2 [[10. 11. 12. 13. 14. 15. 16. 17. 18. 19.]
   [11. 12. 13. 14. 15. 16. 17. 18. 19. 20.]
   [12. 13. 14. 15. 16. 17. 18. 19. 20. 21.]
   [13. 14. 15. 16. 17. 18. 19. 20. 21. 22.]
   [14. 15. 16. 17. 18. 19. 20. 21. 22. 23.]] [29. 30. 31. 32. 33.]
3 [[15. 16. 17. 18. 19. 20. 21. 22. 23. 24.]
   [16. 17. 18. 19. 20. 21. 22. 23. 24. 25.]
   [17. 18. 19. 20. 21. 22. 23. 24. 25. 26.]
   [18. 19. 20. 21. 22. 23. 24. 25. 26. 27.]
   [19. 20. 21. 22. 23. 24. 25. 26. 27. 28.]] [34. 35. 36. 37. 38.]
4 [[20. 21. 22. 23. 24. 25. 26. 27. 28. 29.]
   [21. 22. 23. 24. 25. 26. 27. 28. 29. 30.]
   [22. 23. 24. 25. 26. 27. 28. 29. 30. 31.]
   [23. 24. 25. 26. 27. 28. 29. 30. 31. 32.]
   [24. 25. 26. 27. 28. 29. 30. 31. 32. 33.]] [39. 40. 41. 42. 43.]
5 [[25. 26. 27. 28. 29. 30. 31. 32. 33. 34.]
   [26. 27. 28. 29. 30. 31. 32. 33. 34. 35.]
   [27. 28. 29. 30. 31. 32. 33. 34. 35. 36.]
   [28. 29. 30. 31. 32. 33. 34. 35. 36. 37.]
   [29. 30. 31. 32. 33. 34. 35. 36. 37. 38.]] [44. 45. 46. 47. 48.]
6 [[30. 31. 32. 33. 34. 35. 36. 37. 38. 39.]
   [31. 32. 33. 34. 35. 36. 37. 38. 39. 40.]
   [32. 33. 34. 35. 36. 37. 38. 39. 40. 41.]
   [33. 34. 35. 36. 37. 38. 39. 40. 41. 42.]
   [34. 35. 36. 37. 38. 39. 40. 41. 42. 43.]] [49. 50. 51. 52. 53.]
7 [[35. 36. 37. 38. 39. 40. 41. 42. 43. 44.]
   [36. 37. 38. 39. 40. 41. 42. 43. 44. 45.]
   [37. 38. 39. 40. 41. 42. 43. 44. 45. 46.]
   [38. 39. 40. 41. 42. 43. 44. 45. 46. 47.]
   [39. 40. 41. 42. 43. 44. 45. 46. 47. 48.]] [54. 55. 56. 57. 58.]
8 [[40. 41. 42. 43. 44. 45. 46. 47. 48. 49.]
   [41. 42. 43. 44. 45. 46. 47. 48. 49. 50.]
   [42. 43. 44. 45. 46. 47. 48. 49. 50. 51.]
   [43. 44. 45. 46. 47. 48. 49. 50. 51. 52.]
   [44. 45. 46. 47. 48. 49. 50. 51. 52. 53.]] [59. 60. 61. 62. 63.]
9 [[45. 46. 47. 48. 49. 50. 51. 52. 53. 54.]
   [46. 47. 48. 49. 50. 51. 52. 53. 54. 55.]
   [47. 48. 49. 50. 51. 52. 53. 54. 55. 56.]
   [48. 49. 50. 51. 52. 53. 54. 55. 56. 57.]
   [49. 50. 51. 52. 53. 54. 55. 56. 57. 58.]] [64. 65. 66. 67. 68.]
10 [[ 0.  1.  2.  3.  4.  5.  6.  7.  8.  9.]
     [ 1.  2.  3.  4.  5.  6.  7.  8.  9. 10.]
     [ 2.  3.  4.  5.  6.  7.  8.  9. 10. 11.]
     [ 3.  4.  5.  6.  7.  8.  9. 10. 11. 12.]
     [ 4.  5.  6.  7.  8.  9. 10. 11. 12. 13.]] [19. 20. 21. 22. 23.]

```

Data preparation and configuration of generators

- We selected that an initial sequence of 12000 readings will be used for training and the rest for testing
- Data are normalized with `StandardScaler`

Note in further experiments we will repeat cells initializing the generators. They start at the first index and then they wrap over the end to the beginning.

```
In [67]: x_orig = df.cnt.to_numpy()
from sklearn.preprocessing import StandardScaler
x=x_orig.reshape(-1, 1)
scaler = StandardScaler()
scaler.fit(x[:12_000])
x_scaled=scaler.transform(x)
x_scaled=x_scaled.reshape(-1)
print(x_scaled)
```

```
[-0.94185161 -0.78331641 -0.83616147 ... -0.45303472 -0.6445981
 -0.7238657 ]
```

```
In [68]: window_size = 14*24
horizon = 7*24
batch_size=128
train_gen = generator(x_scaled,
                      window_size=window_size,
                      horizon=horizon,
                      min_index=0,
                      max_index=12_000,
                      shuffle=False,
                      step=1,
                      batch_size=batch_size)
val_gen = generator(x_scaled,
                   window_size=window_size,
                   horizon=horizon,
                   min_index=12_000,
                   max_index=None,
                   shuffle=False,
                   step=1,
                   batch_size=batch_size)

train_steps = (12_000-window_size-horizon-batch_size)//window_size
print(train_steps)

val_steps = (x_scaled.size-12_000-window_size-horizon-batch_size)//window_size
print(val_steps)
```

```
33
14
```

7.4.2 Basic architecture

The basic architecture uses the whole input sequence at once to make the prediction.

```
In [69]: from keras.models import Sequential
from keras import layers
import tensorflow
from tensorflow.keras.optimizers import RMSprop
```



```
import tensorflow as tf

seed=42
np.random.seed(seed)
tf.random.set_seed(seed)

model = Sequential()
model.add(layers.Dense(units=64, input_shape=(window_size,)))
model.add(layers.Dense(units=32, activation='relu'))
model.add(layers.Dense(1))

# model.compile(optimizer=RMSprop(learning_rate=0.0001), loss='mse', met
model.compile(optimizer='adam', loss='mse', metrics=['mse', 'mae'])
model.summary()
```

Model: "sequential_15"

Layer (type)	Output Shape	Param #
dense_29 (Dense)	(None, 64)	21568
dense_30 (Dense)	(None, 32)	2080
dense_31 (Dense)	(None, 1)	33

```
=====  
Total params: 23,681  
Trainable params: 23,681  
Non-trainable params: 0  
=====
```

```
In [70]: hist = model.fit(train_gen, steps_per_epoch=train_steps, epochs=20, batch_si
                                validation_data=val_gen,
                                validation_steps=val_step
```

Epoch 1/20
33/33 [=====] - 1s 14ms/step - loss: 0.7914 - mse: 0.7914 - mae: 0.6913 - val_loss: 2.4575 - val_mse: 2.4575 - val_mae: 1.2735

Epoch 2/20
33/33 [=====] - 0s 1ms/step - loss: 0.8463 - mse: 0.8463 - mae: 0.7089 - val_loss: 8.9027 - val_mse: 8.9027 - val_mae: 2.5366

Epoch 3/20
33/33 [=====] - 0s 1ms/step - loss: 0.9909 - mse: 0.9909 - mae: 0.7843 - val_loss: 2.1246 - val_mse: 2.1246 - val_mae: 1.1339

Epoch 4/20
33/33 [=====] - 0s 1ms/step - loss: 0.5113 - mse: 0.5113 - mae: 0.5427 - val_loss: 5.8752 - val_mse: 5.8752 - val_mae: 1.9840

Epoch 5/20
33/33 [=====] - 0s 1ms/step - loss: 0.6231 - mse: 0.6231 - mae: 0.6078 - val_loss: 3.4087 - val_mse: 3.4087 - val_mae: 1.4633

Epoch 6/20
33/33 [=====] - 0s 1ms/step - loss: 0.6550 - mse: 0.6550 - mae: 0.6073 - val_loss: 4.0740 - val_mse: 4.0740 - val_mae: 1.7074

Epoch 7/20
33/33 [=====] - 0s 1ms/step - loss: 0.5344 - mse: 0.5344 - mae: 0.5540 - val_loss: 2.8495 - val_mse: 2.8495 - val_mae: 1.3932

Epoch 8/20
33/33 [=====] - 0s 1ms/step - loss: 0.5551 - mse: 0.5551 - mae: 0.5373 - val_loss: 0.8861 - val_mse: 0.8861 - val_mae: 0.7238

Epoch 9/20
33/33 [=====] - 0s 1ms/step - loss: 0.3685 - mse: 0.3685 - mae: 0.4675 - val_loss: 2.8953 - val_mse: 2.8953 - val_mae: 1.5320

Epoch 10/20
33/33 [=====] - 0s 1ms/step - loss: 0.4720 - mse: 0.4720 - mae: 0.5027 - val_loss: 3.4896 - val_mse: 3.4896 - val_mae: 1.4977

Epoch 11/20
33/33 [=====] - 0s 1ms/step - loss: 0.5321 - mse: 0.5321 - mae: 0.5407 - val_loss: 3.9955 - val_mse: 3.9955 - val_mae: 1.7239

Epoch 12/20
33/33 [=====] - 0s 1ms/step - loss: 0.3172 - mse: 0.3172 - mae: 0.4263 - val_loss: 6.0127 - val_mse: 6.0127 - val_mae: 2.2152

Epoch 13/20
33/33 [=====] - 0s 1ms/step - loss: 0.5501 - mse: 0.5501 - mae: 0.5442 - val_loss: 4.0822 - val_mse: 4.0822 - val_mae: 1.7028

Epoch 14/20
33/33 [=====] - 0s 1ms/step - loss: 0.6682 - mse: 0.6682 - mae: 0.6300 - val_loss: 4.5407 - val_mse: 4.5407 - val_mae: 1.8822

Epoch 15/20
33/33 [=====] - 0s 1ms/step - loss: 0.4874 - mse: 0.4874 - mae: 0.5150 - val_loss: 4.8723 - val_mse: 4.8723 - val_mae: 1.7895

```

Epoch 16/20
33/33 [=====] - 0s 1ms/step - loss: 0.5393 - ms
e: 0.5393 - mae: 0.5442 - val_loss: 0.8829 - val_mse: 0.8829 - val_mae:
0.7207
Epoch 17/20
33/33 [=====] - 0s 1ms/step - loss: 0.4801 - ms
e: 0.4801 - mae: 0.5208 - val_loss: 4.7497 - val_mse: 4.7497 - val_mae:
2.0368
Epoch 18/20
33/33 [=====] - 0s 2ms/step - loss: 0.4745 - ms
e: 0.4745 - mae: 0.5079 - val_loss: 3.8968 - val_mse: 3.8968 - val_mae:
1.5782
Epoch 19/20
33/33 [=====] - 0s 2ms/step - loss: 0.6814 - ms
e: 0.6814 - mae: 0.6054 - val_loss: 2.4319 - val_mse: 2.4319 - val_mae:
1.2763
Epoch 20/20
33/33 [=====] - 0s 2ms/step - loss: 0.3163 - ms
e: 0.3163 - mae: 0.4251 - val_loss: 3.1546 - val_mse: 3.1546 - val_mae:
1.4452

```

To test on the remaining data, we will configure a test generator, which produces batches containing one single element. The sizes of predicted values `y_pred` and test values `y_test` should match.

```

In [71]: test_gen = generator(x_scaled,
                             window_size=window_size,
                             horizon=horizon,
                             min_index=12_000,
                             max_index=None,
                             shuffle=False,
                             step=1,
                             batch_size=1)

steps = len(x_scaled)-12_000

y_pred = model.predict(test_gen, steps=steps)
print(y_pred.shape)
y_pred = y_pred.reshape(-1)
print(y_pred.shape)

5379/5379 [=====] - 2s 366us/step
(5379, 1)
(5379,)

```

```

In [72]: y_test = x_scaled[12_000:]
print(y_test.shape)
# print(y_test)

(5379,)

```

We will compute some basic regression scores. Analyzing `r2` (determination coefficient) we state that predictive capabilities of our model are not that great.

```

In [73]: import sklearn.metrics

def compute_scores(y_test, y_pred):
    scores={'r2':sklearn.metrics.r2_score,
           'mse':sklearn.metrics.mean_squared_error,

```

```

    'rmse': lambda y_true, y_pred : np.sqrt(sklearn.metrics.mean_squared_error(y_true, y_pred)),
    'maxe': sklearn.metrics.max_error,
    'med': sklearn.metrics.median_absolute_error,
    'mae': sklearn.metrics.mean_absolute_error,
    'mape': sklearn.metrics.mean_absolute_percentage_error,
}
results={}
for k in scores:
    results[k] = scores[k](y_test,y_pred)
return results

r = compute_scores(y_test,y_pred)
print(r)

```

```

{'r2': -0.42897834424032966, 'mse': 3.0113139999171836, 'rmse': 1.7353138044507062, 'maxe': 6.197833985519388, 'med': 1.027468854636469, 'mae': 1.3318156623588446, 'mape': 3.848662361940473}

```

To prepare plots in original units, we must apply inverse transformation to that introduced by the scaler.

```


In [74]: y_test_orig = scaler.inverse_transform(y_test.reshape(-1, 1))
y_pred_orig = scaler.inverse_transform(y_pred.reshape(-1, 1))
y_test_orig = y_test_orig.reshape(-1)
y_pred_orig = y_pred_orig.reshape(-1)
# print(y_test_orig.shape)

```

```

In [75]: plt.plot(y_test_orig[100:200],label='test data')
plt.plot(y_pred_orig[100:200],label='predictions')
plt.legend()
plt.show()

```



7.4.3 GRU

We will apply a basic GRU configuration.

TODO 7.4.1 Make the model more complex, eg. increase a number of units and introduce a small dense layer.

```
In [116... from keras.models import Sequential
from keras import layers
#from keras.optimizers import RMSprop
import tensorflow
from tensorflow.keras.optimizers import RMSprop

model = Sequential()

model.add(layers.GRU(units = 64, input_shape=(None,window_size)))
model.add(layers.Dense(8))
model.add(layers.Dense(1))

model.compile(optimizer=RMSprop(), loss='mse', metrics=['mse', 'mae'])
model.summary()
```

Model: "sequential_40"

Layer (type)	Output Shape	Param #
gru_7 (GRU)	(None, 64)	77184
dense_79 (Dense)	(None, 8)	520
dense_80 (Dense)	(None, 1)	9
Total params: 77,713		
Trainable params: 77,713		
Non-trainable params: 0		

In this case we applied `expand_dims` option - to convert 2D matrix `batch_size x window_size` into `batch_size x window_size x features`. The number of features is 1.

```
In [100... train_gen = generator(x_scaled,
                        window_size=window_size,
                        horizon=horizon,
                        min_index=0,
                        max_index=12_000,
                        shuffle=False,
                        step=1,
                        batch_size=batch_size,
                        expand_dims=True)
val_gen = generator(x_scaled,
                    window_size=window_size,
                    horizon=horizon,
                    min_index=12_000,
                    max_index=None,
                    shuffle=False,
                    step=1,
                    batch_size=batch_size,
                    expand_dims=True)
```

```
train_steps = (12_000-window_size-horizon-batch_size)//window_size
print(train_steps)

val_steps = (x_scaled.size-12_000-window_size-horizon-batch_size)//window
print(val_steps)
```

33

14

In [101... `X,y = next(train_gen)`
`print(X.shape)`

(128, 336, 1)

In [102... `hist = model.fit(train_gen,steps_per_epoch=train_steps,`
`epochs=20,`
`batch_size=batch_size,`
`validation_data=val_gen,`
`validation_steps=val_steps)`

Epoch 1/20
33/33 [=====] - 1s 16ms/step - loss: 0.3556 - mse: 0.2156 - mae: 0.3556 - val_loss: 0.4574 - val_mse: 0.3682 - val_mae: 0.4574

Epoch 2/20
33/33 [=====] - 0s 15ms/step - loss: 0.3869 - mse: 0.3830 - mae: 0.3869 - val_loss: 1.1010 - val_mse: 1.8897 - val_mae: 1.1010

Epoch 3/20
33/33 [=====] - 1s 16ms/step - loss: 0.4512 - mse: 0.4171 - mae: 0.4512 - val_loss: 0.6982 - val_mse: 1.0747 - val_mae: 0.6982

Epoch 4/20
33/33 [=====] - 1s 16ms/step - loss: 0.3813 - mse: 0.3596 - mae: 0.3813 - val_loss: 0.5284 - val_mse: 0.4692 - val_mae: 0.5284

Epoch 5/20
33/33 [=====] - 0s 15ms/step - loss: 0.3569 - mse: 0.2997 - mae: 0.3569 - val_loss: 0.7631 - val_mse: 1.1928 - val_mae: 0.7631

Epoch 6/20
33/33 [=====] - 1s 16ms/step - loss: 0.3990 - mse: 0.3410 - mae: 0.3990 - val_loss: 0.5984 - val_mse: 0.7061 - val_mae: 0.5984

Epoch 7/20
33/33 [=====] - 1s 15ms/step - loss: 0.3418 - mse: 0.3168 - mae: 0.3418 - val_loss: 0.7062 - val_mse: 0.9907 - val_mae: 0.7062

Epoch 8/20
33/33 [=====] - 1s 15ms/step - loss: 0.4178 - mse: 0.3958 - mae: 0.4178 - val_loss: 0.5408 - val_mse: 0.5767 - val_mae: 0.5408

Epoch 9/20
33/33 [=====] - 1s 16ms/step - loss: 0.3101 - mse: 0.2057 - mae: 0.3101 - val_loss: 0.5016 - val_mse: 0.4679 - val_mae: 0.5016

Epoch 10/20
33/33 [=====] - 1s 15ms/step - loss: 0.3616 - mse: 0.3395 - mae: 0.3616 - val_loss: 0.6711 - val_mse: 1.0438 - val_mae: 0.6711

Epoch 11/20
33/33 [=====] - 0s 15ms/step - loss: 0.4159 - mse: 0.3870 - mae: 0.4159 - val_loss: 0.5680 - val_mse: 0.6867 - val_mae: 0.5680

Epoch 12/20
33/33 [=====] - 1s 16ms/step - loss: 0.3136 - mse: 0.2468 - mae: 0.3136 - val_loss: 0.5952 - val_mse: 0.6548 - val_mae: 0.5952

Epoch 13/20
33/33 [=====] - 1s 15ms/step - loss: 0.3769 - mse: 0.3324 - mae: 0.3769 - val_loss: 0.5254 - val_mse: 0.6036 - val_mae: 0.5254

Epoch 14/20
33/33 [=====] - 1s 15ms/step - loss: 0.4065 - mse: 0.3417 - mae: 0.4065 - val_loss: 0.4589 - val_mse: 0.4388 - val_mae: 0.4589

Epoch 15/20
33/33 [=====] - 0s 15ms/step - loss: 0.3294 - mse: 0.2943 - mae: 0.3294 - val_loss: 0.6634 - val_mse: 1.0137 - val_mae: 0.6634


```

Epoch 16/20
33/33 [=====] - 1s 16ms/step - loss: 0.3953 - m
se: 0.3553 - mae: 0.3953 - val_loss: 0.4733 - val_mse: 0.4675 - val_mae:
0.4733
Epoch 17/20
33/33 [=====] - 1s 16ms/step - loss: 0.3465 - m
se: 0.2314 - mae: 0.3465 - val_loss: 0.4671 - val_mse: 0.4098 - val_mae:
0.4671
Epoch 18/20
33/33 [=====] - 0s 15ms/step - loss: 0.3390 - m
se: 0.3125 - mae: 0.3390 - val_loss: 0.6018 - val_mse: 0.9035 - val_mae:
0.6018
Epoch 19/20
33/33 [=====] - 1s 18ms/step - loss: 0.3935 - m
se: 0.3696 - mae: 0.3935 - val_loss: 0.4443 - val_mse: 0.4183 - val_mae:
0.4443
Epoch 20/20
33/33 [=====] - 1s 19ms/step - loss: 0.2700 - m
se: 0.1580 - mae: 0.2700 - val_loss: 0.7127 - val_mse: 1.0537 - val_mae:
0.7127

```

TODO 7.4.2 Display the history content

```
In [ ]: import matplotlib.pyplot as plt
```



TODO 7.4.3 Convert predictions to original range of values and plot an interesting part. The example plots show predictions from the range [100:200]

```
In [81]: y_test_orig = scaler.inverse_transform(y_test.reshape(-1, 1))
y_pred_orig = scaler.inverse_transform(y_pred.reshape(-1, 1))
y_test_orig = y_test_orig.reshape(-1)
y_pred_orig = y_pred_orig.reshape(-1)
```

```
In [82]: r = compute_scores(y_test, y_pred)
print(r)
```




```
{'r2': -0.42897834424032966, 'mse': 3.0113139999171836, 'rmse': 1.735313
8044507062, 'maxe': 6.197833985519388, 'med': 1.027468854636469, 'mae':
1.3318156623588446, 'mape': 3.848662361940473}
```

```
In [83]: r = compute_scores(y_test_orig,y_pred_orig)
print(r)
```

```
{'r2': -0.42897835119704064, 'mse': 69012.39148445425, 'rmse': 262.70209
64599526, 'maxe': 938.2648628354073, 'med': 155.54432678222656, 'mae': 2
01.61815425139153, 'mape': 2.434462716420465}
```

```
In [84]: plt.plot(y_test_orig[100:200],label='test data')
plt.plot(y_pred_orig[100:200],label='predictions')
plt.legend()
plt.show()
```



7.4.4 Stacked GRU

This is rather a time-demanding configuration, therefore we will train it during a few epochs.

```
In [85]: from keras.models import Sequential
from keras import layers
#from keras.optimizers import RMSprop
import tensorflow
from tensorflow.keras.optimizers import RMSprop
from keras.layers import BatchNormalization
from keras import regularizers

model = Sequential()
model.add(layers.GRU(32, dropout=0.05,
                    recurrent_dropout=0.05,
                    return_sequences=True,
                    input_shape=(window_size,1),
                    kernel_regularizer=regularizers.l2(0.001)))
```

```

model.add(layers.GRU(64, activation='relu',
                      dropout=0.05,
                      recurrent_dropout=0.05,
                      kernel_regularizer=regularizers.l2(0.001)))
model.add(layers.Dense(1))

model.compile(optimizer=RMSprop(), loss='mae', metrics=['mse', 'mae'])

```

```

In [103... train_gen = generator(x_scaled,
                             window_size=window_size,
                             horizon=horizon,
                             min_index=0,
                             max_index=12_000,
                             shuffle=False,
                             step=1,
                             batch_size=batch_size,
                             expand_dims=True)
val_gen = generator(x_scaled,
                    window_size=window_size,
                    horizon=horizon,
                    min_index=12_000,
                    max_index=None,
                    shuffle=False,
                    step=1,
                    batch_size=batch_size,
                    expand_dims=True)

train_steps = (12_000-window_size-horizon-batch_size)//window_size
print(train_steps)

val_steps = (x_scaled.size-12_000-window_size-horizon-batch_size)//window_size
print(val_steps)

33
14

```

```

In [104... hist = model.fit(train_gen, steps_per_epoch=train_steps,
                             epochs=5,
                             batch_size=batch_size,
                             validation_data=val_gen,
                             validation_steps=val_steps)

```

```

Epoch 1/5
33/33 [=====] - 1s 17ms/step - loss: 0.2807 - m
se: 0.1635 - mae: 0.2807 - val_loss: 0.4034 - val_mse: 0.3058 - val_mae:
0.4034
Epoch 2/5
33/33 [=====] - 1s 15ms/step - loss: 0.3538 - m
se: 0.3392 - mae: 0.3538 - val_loss: 0.7619 - val_mse: 1.0944 - val_mae:
0.7619
Epoch 3/5
33/33 [=====] - 1s 16ms/step - loss: 0.3894 - m
se: 0.3641 - mae: 0.3894 - val_loss: 0.6422 - val_mse: 0.9231 - val_mae:
0.6422
Epoch 4/5
33/33 [=====] - 1s 15ms/step - loss: 0.3143 - m
se: 0.2698 - mae: 0.3143 - val_loss: 0.4538 - val_mse: 0.3735 - val_mae:
0.4538
Epoch 5/5
33/33 [=====] - 0s 15ms/step - loss: 0.3301 - m
se: 0.2692 - mae: 0.3301 - val_loss: 0.6402 - val_mse: 0.9436 - val_mae:
0.6402

```

TODO 7.4.3 Testing with `batch_size=1` is rather a long process.

- Configure `test_gen` to produce larger batches.
- Compute a number of steps to be done to cover testing data [12000:]
- Some data at the end will not be tested
- Select appropriate subsequence of `x_scaled`
- Compute scores

```

In [88]: length = len(x_scaled)-12_000

y_pred = model.predict(test_gen,steps=length)
y_pred = y_pred.reshape(-1)
print(y_pred.shape)

5379/5379 [=====] - 96s 18ms/step
(5379,)

In [89]: y_test = x_scaled[12_000:]
print(y_test)

[-0.99469668 -0.88240091 -0.45303472 ... -0.45303472 -0.6445981
-0.7238657 ]

```

TODO 7.4.4 Return to the original ranges and plot the data

```


In [90]: y_test_orig = scaler.inverse_transform(y_test.reshape(-1, 1))
y_pred_orig = scaler.inverse_transform(y_pred.reshape(-1, 1))
y_test_orig = y_test_orig.reshape(-1)
y_pred_orig = y_pred_orig.reshape(-1)

r = compute_scores(y_test_orig,y_pred_orig)
print(r)

{'r2': -0.09032459280500627, 'mse': 52657.136184571675, 'rmse': 229.4714
2781743366, 'maxe': 863.8881225585938, 'med': 129.95213317871094, 'mae':
171.84590732601342, 'mape': 3.53883617286483}

```

```
In [91]: plt.plot(y_test_orig[100:200], label='test data')
plt.plot(y_pred_orig[100:200], label='predictions')
plt.legend()
plt.show()
```



7.4.5 Conv1D - applying 1D convolution

Finally, we will perform experiment with 1D convolutional filters

```
In [121... from keras.models import Sequential
from keras import layers
model = Sequential()
model.add(layers.Conv1D(filters=32, kernel_size=7, input_shape=(window_si
model.add(layers.MaxPooling1D(pool_size=2))
model.add(layers.Conv1D(filters=64, kernel_size=7, activation='relu', padd
model.add(layers.MaxPooling1D(pool_size=2))
model.add(layers.Flatten())
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1))

model.compile(optimizer='adam', loss='mae', metrics=['mse', 'mae'])
model.summary()
```

Model: "sequential_41"

Layer (type)	Output Shape	Param #
conv1d_6 (Conv1D)	(None, 336, 32)	256
max_pooling1d_4 (MaxPooling1D)	(None, 168, 32)	0
conv1d_7 (Conv1D)	(None, 168, 64)	14400
max_pooling1d_5 (MaxPooling1D)	(None, 84, 64)	0
flatten_1 (Flatten)	(None, 5376)	0
dense_81 (Dense)	(None, 16)	86032
dense_82 (Dense)	(None, 1)	17

```

=====
Total params: 100,705
Trainable params: 100,705
Non-trainable params: 0
=====

```

```

In [119]: train_gen = generator(x_scaled,
                                window_size=window_size,
                                horizon=horizon,
                                min_index=0,
                                max_index=12_000,
                                shuffle=False,
                                step=1,
                                batch_size=batch_size,
                                expand_dims=True)
val_gen = generator(x_scaled,
                    window_size=window_size,
                    horizon=horizon,
                    min_index=12_000,
                    max_index=None,
                    shuffle=False,
                    step=1,
                    batch_size=batch_size,
                    expand_dims=True)
train_steps = (12_000-window_size-horizon-batch_size)//window_size
print(train_steps)

val_steps = (x_scaled.size-12_000-window_size-horizon-batch_size)//window_size
print(val_steps)

33
14

```

```

In [94]: hist = model.fit(train_gen, steps_per_epoch=train_steps,
                           epochs=20,
                           batch_size=batch_size,
                           validation_data=val_gen,
                           validation_steps=val_steps)

```

Epoch 1/20
33/33 [=====] - 1s 18ms/step - loss: 0.5885 - mse: 0.7080 - mae: 0.5885 - val_loss: 1.1053 - val_mse: 2.5588 - val_mae: 1.1053

Epoch 2/20
33/33 [=====] - 1s 16ms/step - loss: 0.5550 - mse: 0.6699 - mae: 0.5550 - val_loss: 1.9004 - val_mse: 4.8980 - val_mae: 1.9004

Epoch 3/20
33/33 [=====] - 0s 14ms/step - loss: 0.7142 - mse: 1.0982 - mae: 0.7142 - val_loss: 1.0366 - val_mse: 1.9365 - val_mae: 1.0366

Epoch 4/20
33/33 [=====] - 0s 14ms/step - loss: 0.6023 - mse: 0.7853 - mae: 0.6023 - val_loss: 1.3446 - val_mse: 3.1480 - val_mae: 1.3446

Epoch 5/20
33/33 [=====] - 0s 14ms/step - loss: 0.4771 - mse: 0.5225 - mae: 0.4771 - val_loss: 1.1466 - val_mse: 2.4045 - val_mae: 1.1466

Epoch 6/20
33/33 [=====] - 0s 14ms/step - loss: 0.7196 - mse: 1.0789 - mae: 0.7196 - val_loss: 1.1350 - val_mse: 2.3973 - val_mae: 1.1350

Epoch 7/20
33/33 [=====] - 0s 14ms/step - loss: 0.5847 - mse: 0.7957 - mae: 0.5847 - val_loss: 1.3251 - val_mse: 2.9392 - val_mae: 1.3251

Epoch 8/20
33/33 [=====] - 0s 14ms/step - loss: 0.5040 - mse: 0.5900 - mae: 0.5040 - val_loss: 0.7027 - val_mse: 0.8237 - val_mae: 0.7027

Epoch 9/20
33/33 [=====] - 0s 14ms/step - loss: 0.4911 - mse: 0.4556 - mae: 0.4911 - val_loss: 0.6437 - val_mse: 0.8418 - val_mae: 0.6437

Epoch 10/20
33/33 [=====] - 0s 14ms/step - loss: 0.4653 - mse: 0.4770 - mae: 0.4653 - val_loss: 1.0951 - val_mse: 2.1039 - val_mae: 1.0951

Epoch 11/20
33/33 [=====] - 0s 14ms/step - loss: 0.5395 - mse: 0.5699 - mae: 0.5395 - val_loss: 0.6826 - val_mse: 0.8314 - val_mae: 0.6826

Epoch 12/20
33/33 [=====] - 0s 14ms/step - loss: 0.3683 - mse: 0.2772 - mae: 0.3683 - val_loss: 0.7809 - val_mse: 1.0017 - val_mae: 0.7809

Epoch 13/20
33/33 [=====] - 0s 15ms/step - loss: 0.4490 - mse: 0.4510 - mae: 0.4490 - val_loss: 0.7990 - val_mse: 1.0544 - val_mae: 0.7990

Epoch 14/20
33/33 [=====] - 0s 14ms/step - loss: 0.4777 - mse: 0.4764 - mae: 0.4777 - val_loss: 0.8238 - val_mse: 1.0878 - val_mae: 0.8238

Epoch 15/20
33/33 [=====] - 0s 14ms/step - loss: 0.4065 - mse: 0.3846 - mae: 0.4065 - val_loss: 0.8260 - val_mse: 1.2973 - val_mae: 0.8260

```

Epoch 16/20
33/33 [=====] - 0s 14ms/step - loss: 0.4133 - m
se: 0.4003 - mae: 0.4133 - val_loss: 0.6453 - val_mse: 0.6994 - val_mae:
0.6453
Epoch 17/20
33/33 [=====] - 0s 14ms/step - loss: 0.3890 - m
se: 0.3004 - mae: 0.3890 - val_loss: 0.5561 - val_mse: 0.6058 - val_mae:
0.5561
Epoch 18/20
33/33 [=====] - 0s 14ms/step - loss: 0.3748 - m
se: 0.3563 - mae: 0.3748 - val_loss: 0.7097 - val_mse: 1.0871 - val_mae:
0.7097
Epoch 19/20
33/33 [=====] - 0s 14ms/step - loss: 0.4478 - m
se: 0.4423 - mae: 0.4478 - val_loss: 0.5020 - val_mse: 0.5068 - val_mae:
0.5020
Epoch 20/20
33/33 [=====] - 0s 14ms/step - loss: 0.3291 - m
se: 0.2233 - mae: 0.3291 - val_loss: 0.8171 - val_mse: 1.2689 - val_mae:
0.8171

```

TODO 7.4.5 Make predictions, compute regression scores

```

In [122]: length = len(x_scaled)-12_000

y_pred = model.predict(test_gen,steps=length)
print(y_pred.shape)
y_pred = y_pred.reshape(-1)
print(y_pred.shape)

5379/5379 [=====] - 3s 449us/step
(5379, 1)
(5379,)

In [96]: y_test = x_scaled[12_000:]
print(y_test.shape)

(5379,)

In [97]: r = compute_scores(y_test,y_pred)
print(r)

{'r2': -0.6729409455981648, 'mse': 3.5254211589837654, 'rmse': 1.8776104
918176628, 'maxe': 5.678658415124616, 'med': 1.2511941096346928, 'mae':
1.516513884616032, 'mape': 3.7080286270421827}

```

TODO 7.4.6 Plot the data returning to the original data ranges

```

In [98]: y_test_orig = scaler.inverse_transform(y_test.reshape(-1, 1))
y_pred_orig = scaler.inverse_transform(y_pred.reshape(-1, 1))
y_test_orig = y_test_orig.reshape(-1)
y_pred_orig = y_pred_orig.reshape(-1)

r = compute_scores(y_test_orig,y_pred_orig)
print(r)

{'r2': -0.6729409267719078, 'mse': 80794.54392855865, 'rmse': 284.243810
71284324, 'maxe': 859.6689910888672, 'med': 189.4131851196289, 'mae': 22
9.5788647138534, 'mape': 6.4519597031302025}

```

```
In [99]: plt.plot(y_test_orig[100:200],label='test data')  
plt.plot(y_pred_orig[100:200],label='predictions')  
plt.legend()  
plt.show()
```

